



THE UNIVERSITY OF
CHICAGO | **Research
Computing
Center**

Slurm Resource Scheduler and Other Best Practices for Using the Computing Cluster

Jonathan Skone

Today's Agenda

SLURM Job Submission on the HPC Cluster

Objectives:

- Be comfortable submitting jobs on the RCC HPC Cluster
- Understand queue priority
- Ability to query status of job and understand exit codes
- Know how to fix common errors
- Using array jobs and GNU parallel

Using DASK on RCC resources (led by Ivan)



THE UNIVERSITY OF
CHICAGO

Office of Research and
National Laboratories
Research Computing Center

Connecting to Midway2 and Setup

Midway2 logins:

midway2.rcc.uchicago.edu

Midway 2

- Mac/Linux users/Windows Powershell users:

Open a terminal app and issue the following command.



```
w/o X11 [johnnyb@volpe]$ ssh $USER@midway2.rcc.uchicago.edu
```

Please see the following link for additional instructions on setting up your client and connecting to Midway2 resources.

[Connecting to midway2 Instructions](#)

- Get the materials for today's workshop:

After logging in to midway2 login nodes. Use the copy 'cp' command to copy today's materials to your scratch directory \$SCRATCH

```
[johnnyb@m21]$ cd $SCRATCH ; cp /project2/env_bootcamp/materials/jobs ./
```



THE UNIVERSITY OF
CHICAGO

Office of Research and
National Laboratories
Research Computing Center

Remember to Respect the Login Nodes

There are more than 1300+ compute nodes that are part of midway2, but there are only 2 login nodes that serve about 300 active users.

There is a process that monitors the login nodes and will kick out users who are using the logins inappropriately.

.Please adhere to the following practices for login node use

- login nodes are for editing files, compiling, moving files, changing permissions, and other non-intensive tasks.
- We recommend to use *sinteractive* for interactive runs
- For long running jobs => submit them to the queue



Review of the RCC compute cluster

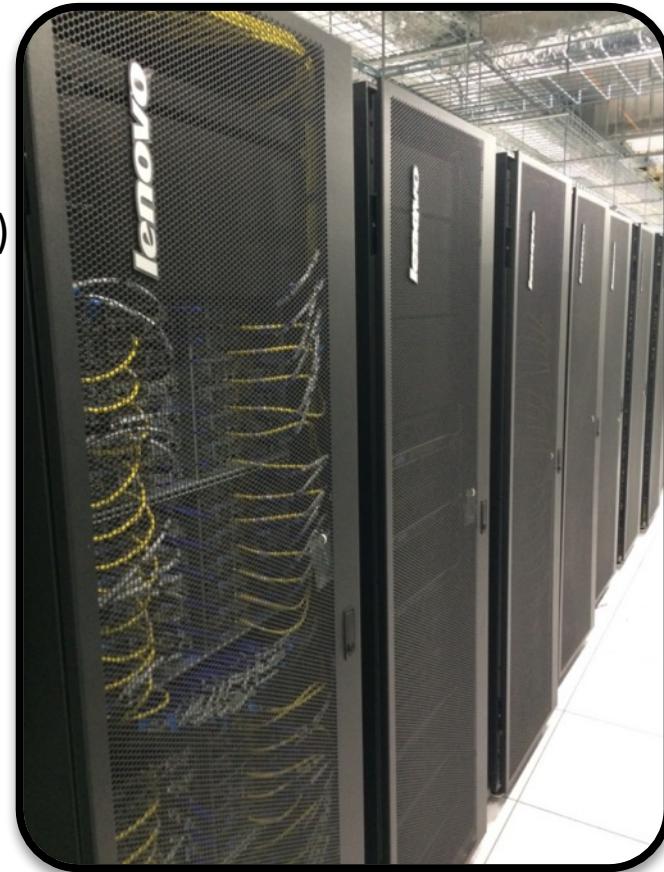
Computing hardware

367 nodes total

- **342** tightly coupled Broadwell nodes (10,360 cores)
 - Two intel E5-2680v4 processors per node (14 core/proc)
 - 155 nodes have EDR network card
 - 187 nodes have FDR network card
- **6** NVidia Tesla K80 GPU nodes (4 GPU cards/node)
- **5** large shared memory nodes (512GB each)
- **14** dual socket loosely-coupled Broadwell nodes

Cluster Partnership Program: 1000+ nodes

- **900+** tightly coupled infiniband nodes
- **20+** Big memory nodes
- **100+** Nvidia GPU nodes



**RCC Manages 1515 nodes
(35,904 cores)**



THE UNIVERSITY OF
CHICAGO

Office of Research and
National Laboratories
Research Computing Center

RCC Midway2 Shared Hardware

To query the list of communal resources that users can submit jobs to use the `rcchelp` command.

```
[johnnyb@m21]$ rcchelp sinfo shared
```

Cores(Nodes)	CPU	Memory	Slurm Queue
5180 (185)	2.4GHz Intel Broadwell	64GB	broadwl
5180 (185)	2.4GHz Intel Broadwell	64GB	broadwl
168 (6)	2.4GHz Intel Broadwell	512GB	bigmem2
504 (18)	2.4GHz Intel Broadwell	64GB	broadwl-lc

Cores (Nodes)	CPU	Memory	Queue
	Accelerator		
168 (6)	2.4GHz Intel Broadwell	128GB	gpu2
	NVIDIA Tesla K80		

NOTE: For this bootcamp you only have access to the `edu` partition



THE UNIVERSITY OF
CHICAGO

Office of Research and
National Laboratories
Research Computing Center

RCC Job Scheduler: SLURM



Open source and very powerful job scheduler.

Most HPC centers have migrated to using This scheduler.

We are currently running slurm v. 18.08.8



SLURM documentation

<https://slurm.schedmd.com/sbatch.html>



THE UNIVERSITY OF
CHICAGO

Office of Research and
National Laboratories
Research Computing Center

Review: Running Jobs Interactively

- Accessing a compute node interactively will allow you to run the job at the command prompt directly on the compute node.
 - To run interactively use the sinteractive command
 - Uses Slurm to provide access to dedicated node(s) to which you can login directly
 - To use sinteractive:
- ```
sinteractive --time=00:30:00 --nodes=1
--ntasks=2 --mem-per-cpu=1000
--partition=edu
```
- **Note:** If you close your terminal you will NOT kill the interactive session. To exit the interactive session, type “exit” from the compute node or use **scancel**



# Slurm: Some key terms to remember

When we refer to a **job** we mean a task or set of tasks that you have you have scheduled to use the compute resources through the SLURM scheduler.

The **queue** is all jobs in a particular partition. They can be in the RUNNING, PENDING, DEPENDENCY, COMPLETING, or HELD state. To see every job in a particular partitions queue, use the command `squeue -p <partition>`

```
[johnnyb@m21]$ squeue -p edu
```

To see only your jobs in the queue use `squeue -u` or `myq`

```
[johnnyb@m21]$ squeue -u $USER
```

```
[johnnyb@m21]$ myq
```



# What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash

Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

# What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash

Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

This **#!** is a shebang

It tells operating system to use  
**/bin/bash**  
with this script

# What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash

Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

# is a comment  
everything after # is ignored by  
bash

# What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash

Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

**#SBATCH** is a directive

It is a comment in Bash

**#SBATCH** is only relevant to slurm:  
sbatch my\_script.sh

# What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash

Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

**#SBATCH** is a directive

It is a comment in Bash

**#SBATCH** is only relevant to slurm:  
sbatch my\_script.sh

To comment out directives, break  
the pattern, e.g.  
##SBATCH  
#SBATCH

# What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash

Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

Instructions for Slurm must go at the top of the script

Any #SBATCH lines you put after your program will be ignored

# What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash

Here is a comment
#SBATCH --time=1:00:00

#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
```

Slurm has some variables you can use. %j is the job number. When the job runs %j will be expanded to the job number. In this example %j is used in the output file and error file names:

MyJob-**13571056**.out  
MyJob-**13571056**.err

**%j** is unique. By using %j in your filenames you guarantee a unique file name, which means you won't accidentally overwrite previous output.

# What goes into a batch script?

A batch script is list of instructions for slurm.

```
#!/bin/bash
Here is a comment
#SBATCH --time=1:00:00 => Time your job is allowed to run

#SBATCH --nodes=1 => Number of nodes to run on
#SBATCH --ntasks-per-node=1 => Number of cores on each node to use
#SBATCH --mem-per-cpu=2000 => Memory per cpu => 2000Mb or 2Gb
#SBATCH --job-name=MyJob => Name of the job.
#SBATCH --output= MyJob-%j.out => Job output file behaves as stdout for the code.
#SBATCH --error=MyJob-%j.err => Error file. behaves as stderr for the code.
#SBATCH --mail-user=jb@uchicago.edu => User to mail. Email address to send job notices
#SBATCH --mail-user=jb@uchicago.edu => Type of email notices to be sent to user.
#SBATCH --mail-type=ALL

module load <module name> => Load any modules you need for your application
python myjob.py => run the code you want
```

# Running batch jobs using a Submission Script

- A simple job submission script (saved as job.sbatch):

```
#!/bin/bash
#SBATCH --job-name=first_python_job
#SBATCH --output=first_python_job_%j.out
#SBATCH --error=first_python_job_%j.err
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000M
#SBATCH --partition=broadwl
#SBATCH --reservation=kicpworkshop-cpu
#SBATCH --time=00:30:00
module load python
python hello_world.py
echo "job finished at `date`"
```

Slurm      Your Job

- To submit the above script:
  - **sbatch** job.sbatch



# Please do Exercise 1 in the job folder

```
[johnnyb@m21]$ cd $SCRATCH
[johnnyb@m21]$ cd jobs/Ex-1
[johnnyb@m21]$ cat job.sbatch
[johnnyb@m21]$ sbatch job.sbatch
[johnnyb@m21]$ squeue -u $USER
```



# How to submit OpenMP jobs

```
#!/bin/bash

#Here is a comment
#SBATCH --time=1:00:00
#SBATCH --partition=broadwl
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=8
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
export OMP_NUM_THREADS=8
#Run your code
./my_executable
```

Applicable to using any python package that relies on threaded BLAS routines coming from Intel MKL

Specify number of cores > 1.

OMP\_NUM\_THREADS is an environment variable.

# Please do Exercise 2 in the job folder

```
[johnnyb@m21]$ cd $SCRATCH
[johnnyb@m21]$ cd jobs/Ex-2
[johnnyb@m21]$ cat job.sbatch
[johnnyb@m21]$ sbatch job.sbatch
[johnnyb@m21]$ squeue -u $USER
```

Note: here we use a Makefile to manage building the program. One can create tasks that are programmatically called using a Makefile



# How to submit Parallel MPI jobs

```
#!/bin/bash

#Here is a comment
#SBATCH --time=1:00:00
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err
#SBATCH --partition=edu
#SBATCH --nodes=4
#SBATCH --ntasks-per-node=8
#SBATCH --mem-per-cpu=2000
module load openmpi
module load <module name>
#Run your code
mpirun ./my_executable
```

Always use –ntasks-per-node for MPI jobs  
In combination w/ --nodes flag

Specify number of nodes > 1.

Specify number of cores >= 1.

Load OPENMPI MPI library or IntelMPI

# Please do Exercise 3 in the job folder

```
[johnnyb@m21]$ cd $SCRATCH
[johnnyb@m21]$ cd jobs/Ex-3
[johnnyb@m21]$ cat job.sbatch
[johnnyb@m21]$ sbatch job.sbatch
[johnnyb@m21]$ squeue -u $USER
```

Note: here we use a Makefile to manage building the program. One can create tasks that are programmatically called using a Makefile



# How to submit GPU jobs?

```
#!/bin/bash

#Here is a comment
#SBATCH --time=1:00:00
#SBATCH --partition=gpu2
#SBATCH --gres=gpu:1
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=8
#SBATCH --mem-per-cpu=2000
#SBATCH --job-name=MyJob
#SBATCH --output= MyJob-%j.out
#SBATCH --error=MyJob-%j.err

module load <module name>
#Run your code
./my_executable
```

Specify partition gpu2

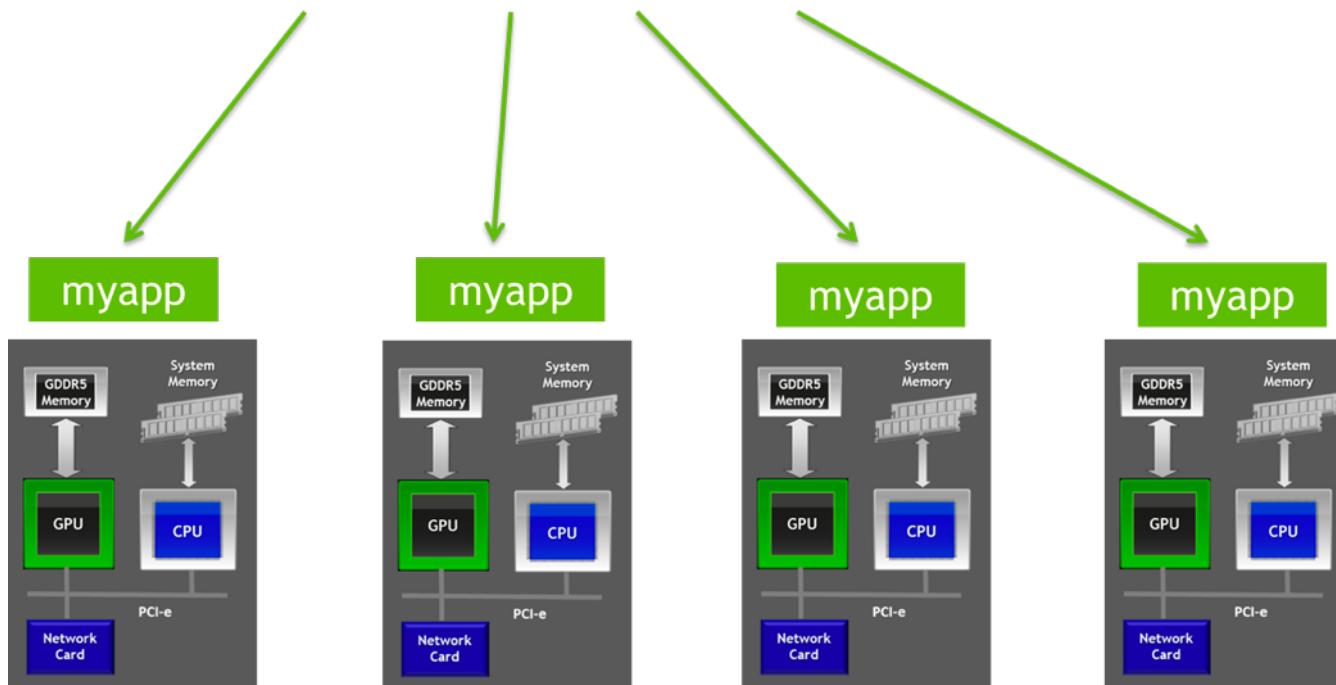
Specify number of gpus, like *gpu:1*

We will not run this example since  
the env\_bootcamp does not have  
access to the gpu2 partition

# How to submit MPI + GPU jobs?

CUDA Aware MPI

```
mpirun -np 4 ./myapp <args>
```



# How to submit MPI + GPU jobs?

## Compilation

```
#!/bin/bash .brown.edu/oscar/gpu-computing/mpi-cuda
SUBMITTING JOBS
Running Jobs The CUDA/C++ comp...
module load openmpi/3.1.2
module load cuda/10.1 Animations Slide Show Review View Te...
Compiling the device code
nvcc -c dev.cu
#Compiling the host code
mpicc -c hostname.c
Ex-4 in Repo

Linking the host and device code
mpicc -o HostMap dev.o hostname.o -lcudart
Ex-4 in Repo

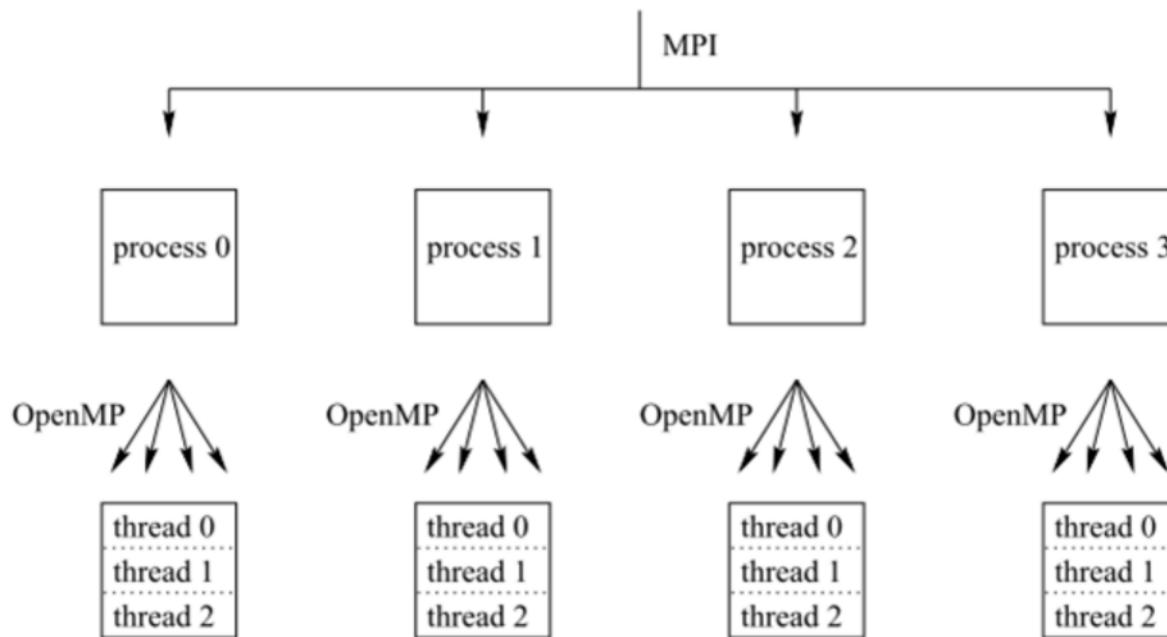
How to submit MPI + GPU jobs?
#Submitting the job as batch script
sbatch mpijob.sh
```

## Job Submission

```
#!/bin/bash .brown.edu/oscar/gpu-computing/mpi-cuda
SUBMITTING JOBS
Running Jobs The C...
#SBATCH -t 00:30:00
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=4
#SBATCH --partition=gpu2
#SBATCH --gres=gpu:2
#SBATCH --job-name=MyJob
#SBATCH --output=MyJob-%j.out
#SBATCH --error=MyJob-%j.err
#SBATCH --qos=stafftest
mpirun ./HostMap
```

See Ex-5 for an example. Note we can not run this in this bootcamp

# How to submit MPI + OpenMPI jobs?



```
#!/bin/bash
#SBATCH --job-name=hybrid
#SBATCH --output=hybrid_%j.out
#SBATCH --error=hybrid_%j.err
#SBATCH --time=00:10:00
#SBATCH --ntasks=4
#SBATCH --cpus-per-task=8
#SBATCH --partition=broadwl
#SBATCH --constraint=edr
#SBATCH --qos=stafftest

Load the default OpenMPI module.
module load openmpi

Set OMP_NUM_THREADS to the number of CPUs per task we asked for.
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

mpirun ./mpomp
```

# How to submit array based jobs

```
#!/bin/bash
Job Name
#SBATCH --job-name=arrayjob
Walltime requested
#SBATCH --time=0:10:00
#Add partition
#SBATCH --partition=edu

Provide index values (TASK IDs)
#SBATCH --array=1-16

Use '%A' for array-job ID, '%J' for job ID and '%a' for task ID
#SBATCH --error=maths%A-%a.err
#SBATCH --output=maths%A-%a.out

single core
#SBATCH --ntasks-per-node=1
#SBATCH --mem-per-cpu=2000

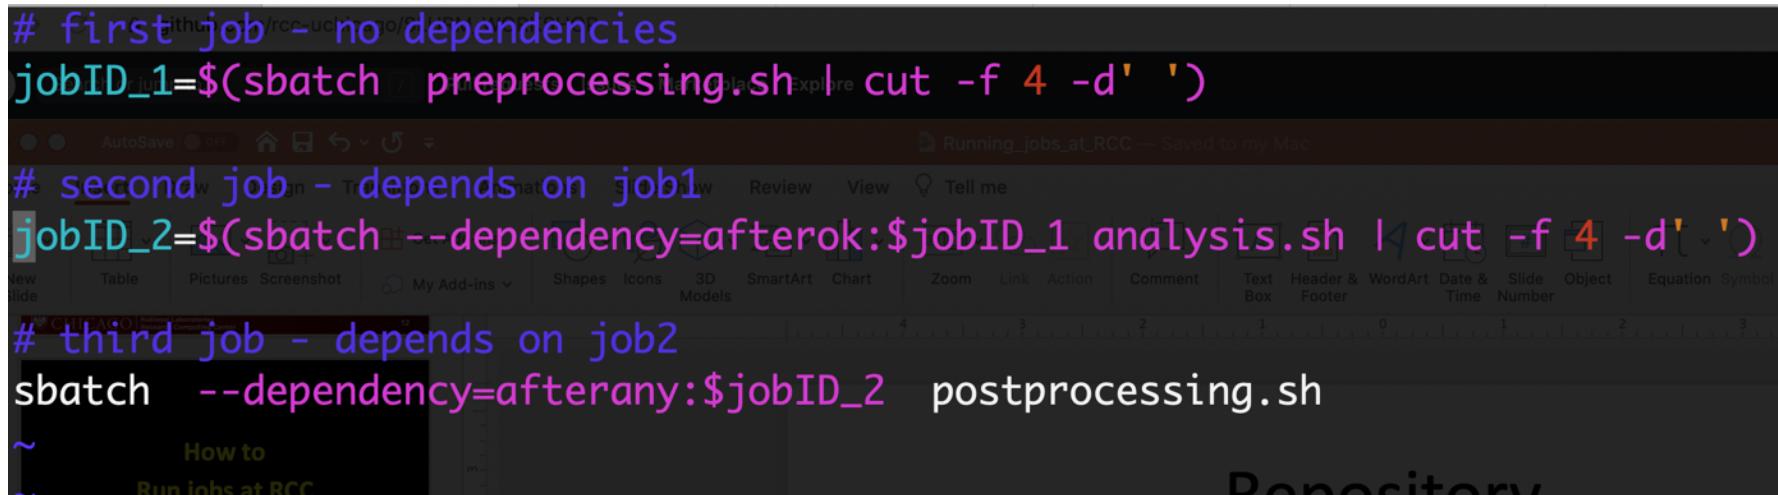
Use the $SLURM_ARRAY_TASK_ID variable to provide different inputs for each job
input=$((SLURM_ARRAY_TASK_ID*1000+2))
echo "Running job array number: "$SLURM_ARRAY_TASK_ID "input " $input
```

Please do the Exercise Ex-6

# How to submit dependent jobs

SLURM Rule:

*sbatch --dependency=type:job\_id jobfile*



A screenshot of a Microsoft Word document titled "Running\_jobs\_at\_RCC". The document contains the following text:

```
first job - no dependencies
jobID_1=$(sbatch preprocessing.sh | cut -f 4 -d' ')

second job - depends on job1
jobID_2=$(sbatch --dependency=afterok:$jobID_1 analysis.sh | cut -f 4 -d' ')

third job - depends on job2
sbatch --dependency=afterany:$jobID_2 postprocessing.sh
```

The code demonstrates how to submit three jobs sequentially using SLURM's dependency features. The first job runs "preprocessing.sh" and outputs its job ID to the command line. The second job runs "analysis.sh" and depends on the completion of the first job. The third job runs "postprocessing.sh" and depends on the completion of either the first or second job.

# How to submit dependent jobs

|            |                                                                                                                              |
|------------|------------------------------------------------------------------------------------------------------------------------------|
| after      | This job can begin execution after the specified jobs have begun execution                                                   |
| afterany   | This job can begin execution after the specified jobs have terminated.                                                       |
| aftercorr  | A task of this job array can begin execution after the corresponding task ID in the specified job has completed successfully |
| afternotok | This job can begin execution after the specified jobs have terminated in some failed state                                   |
| afterok    | This job can begin execution after the specified jobs have successfully executed                                             |
| singleton  | This job can begin execution after any previously launched jobs sharing the same job name and user have terminated           |

Please do the Exercise Ex-7

# How to submit Parallel batch jobs

```
#!/bin/sh

#SBATCH --time=01:00:00
#SBATCH --partition=edu
#SBATCH --ntasks=20
#SBATCH --mem-per-cpu=2G # NOTE DO NOT USE THE --mem= OPTION
Load the default version of GNU parallel. module load parallel
set the max number of processes (which determine the max per processor)
ulimit -u 10000

This specifies the options used to run srun. The "-N1 -n1" options are
used to allocates a single core to each task.

srun="srun --exclusive -N1 -n1"

#Run GNU parallel

parallel="parallel --delay 0.2 -j $SLURM_NTASKS --joblog runtask.log --resume"
Run a script, runtask.sh, using GNU parallel and srun.

$parallel "$srun ./runtask.sh arg1:{1} > runtask.sh.{1}" ::: {1..128}
Note that if your program does not take any input, use the -n0 option to call the
parallel command: ## $parallel -n0

"$srun ./run_noinput_task.sh > output.{1}" ::: {1..128}
```

# How to submit Parallel batch jobs

Please do the Exercise Ex-8

# What resources should I ask for?

This depends on the code you are running

# What resources should I ask for?

This depends on the code you are running

# What resources should I ask for?

This depends on the code you are running

## Nodes/Cores

- Question: is your code parallel? You will need to find out if your code can
  - Run on multiple cores? Run across multiple nodes?
  - Check if your code is threaded, multiprocessor, MPI
- Question: Is your code serial?
  - This means it can only make use of one core

# What resources should I ask for?

This depends on the code you are running

## Wall Time

Make an estimate of your job run and add a bit.

e.g. if think your code will take an hour, give it 1 hour and 30 min

- If your job runs out of time, your job will be killed, so
- be accurate with your estimate without going below.

# What resources should I ask for?

This depends on the code you are running

## Memory

For memory, this can take some trial and error. You can ask for a lot, then measure your usage. If you have asked for more memory and then reduce your memory with the next job.

- To ask for all the memory available on a node, use `#SBATCH --mem=0`

# What if I need an entire node or specific features?

Add this in your batch script

```
#SBATCH --exclusive
```

Add this in your batch script

```
#SBATCH --constraint=v100
```

# How do I know the features of the node to use with #SBATCH -constraint?

## nodestatus

A screenshot of a terminal window showing the output of the `nodestatus` command. The terminal prompt is `[rajshukla@midway2-Login1 ~]$`. The output displays a table of nodes with their details. A white arrow points from the word "Features" in a pink font at the top right towards the "Features" column in the table. The table has columns: NODES, NAME, CORES, FEATURES, STATUS, CORES IN USE, and PURPOSE. The "FEATURES" column lists various hardware specifications for each node.

| NODES        | NAME    | CORES | FEATURES                             | STATUS | CORES IN USE | PURPOSE |
|--------------|---------|-------|--------------------------------------|--------|--------------|---------|
| midway2-0002 | 28-core | 58GB  | tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b | mix    | 16           | 57.1%   |
| midway2-0003 | 28-core | 58GB  | tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b | alloc  | 28           | 100%    |
| midway2-0004 | 28-core | 58GB  | tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b | alloc  | 28           | 100%    |
| midway2-0005 | 28-core | 58GB  | tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b | alloc  | 28           | 100%    |
| midway2-0006 | 28-core | 58GB  | tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b | mix    | 25           | 89.2%   |
| midway2-0007 | 28-core | 58GB  | tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b | mix    | 16           | 57.1%   |
| midway2-0008 | 28-core | 58GB  | tc,e5-2680v4,64GB,ib,fdr,ibspine-d9b | mix    | 19           | 67.8%   |

# What resources should I ask for?

This depends on the code you are running

**GPUs** If your code is built to use GPUs you can submit to the gpu partition. To request 1 GPU:

```
#SBATCH -p gpu2 --gres=gpu:1
```

# What resources did my job actually use?

It is good practice to occasionally check what resources your job is using. For example if you are going to be submitting hundreds of similar jobs, you may save yourself a lot of waiting time in the queue by checking that you are not over requesting resources. Midway2 has a script to display the resources a job used:

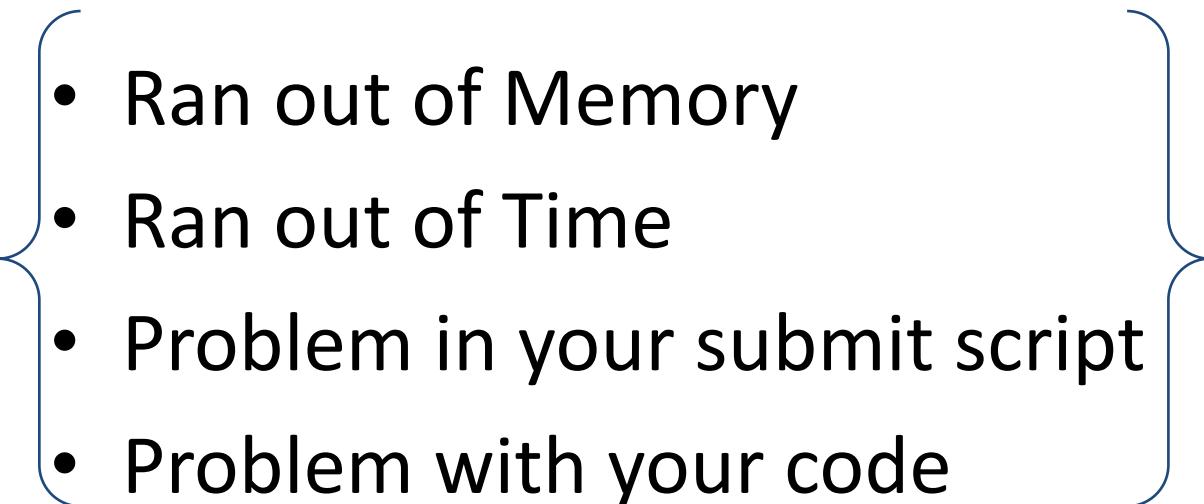
```
jobinfo -j 99999
```

**Replace 99999 with the job ID of the job you are interested in.**

# Why did my job fail?

- Ran out of Memory
- Ran out of Time
- Problem in your submit script
- Problem with your code
- Node failure

# Why did my job fail?

- Ran out of Memory
  - Ran out of Time
  - Problem in your submit script
  - Problem with your code
  - Node failure
- 
- You can fix these

# Why did my job fail?

- Ran out of Memory
  - Ran out of Time
  - Problem in your submit script
  - Problem with your code
  - Node failure
- You can fix these

=> For node failure – please email  
[help@rcc.uchicago.edu](mailto:help@rcc.uchicago.edu), if this happens

# Job Priority

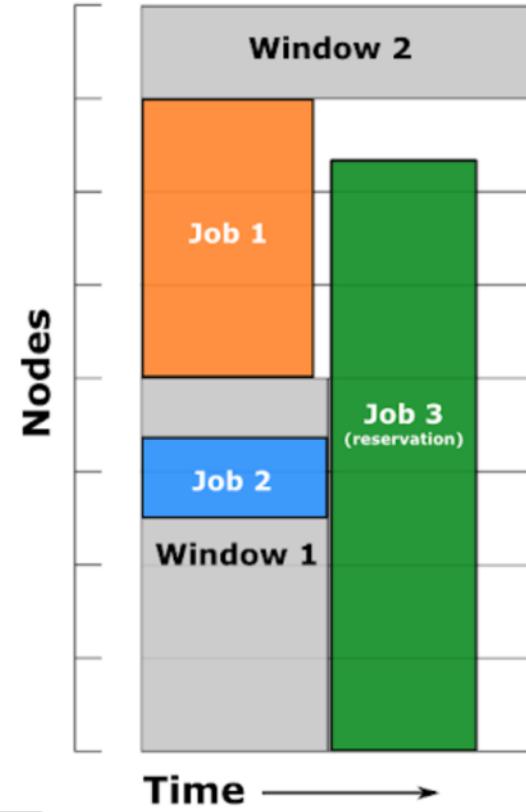
- Priority is calculated using a **Fairshare** algorithm
- Fairshare is function of
  - Requested wall clock, memory, nodes/cores, etc.
  - Length of time in queue
  - Number of jobs in a time window and per PI group
  - Backfill
  - Etc.

# Job Priority

## Backfill



This blue line  
represents all the  
cores on Midway



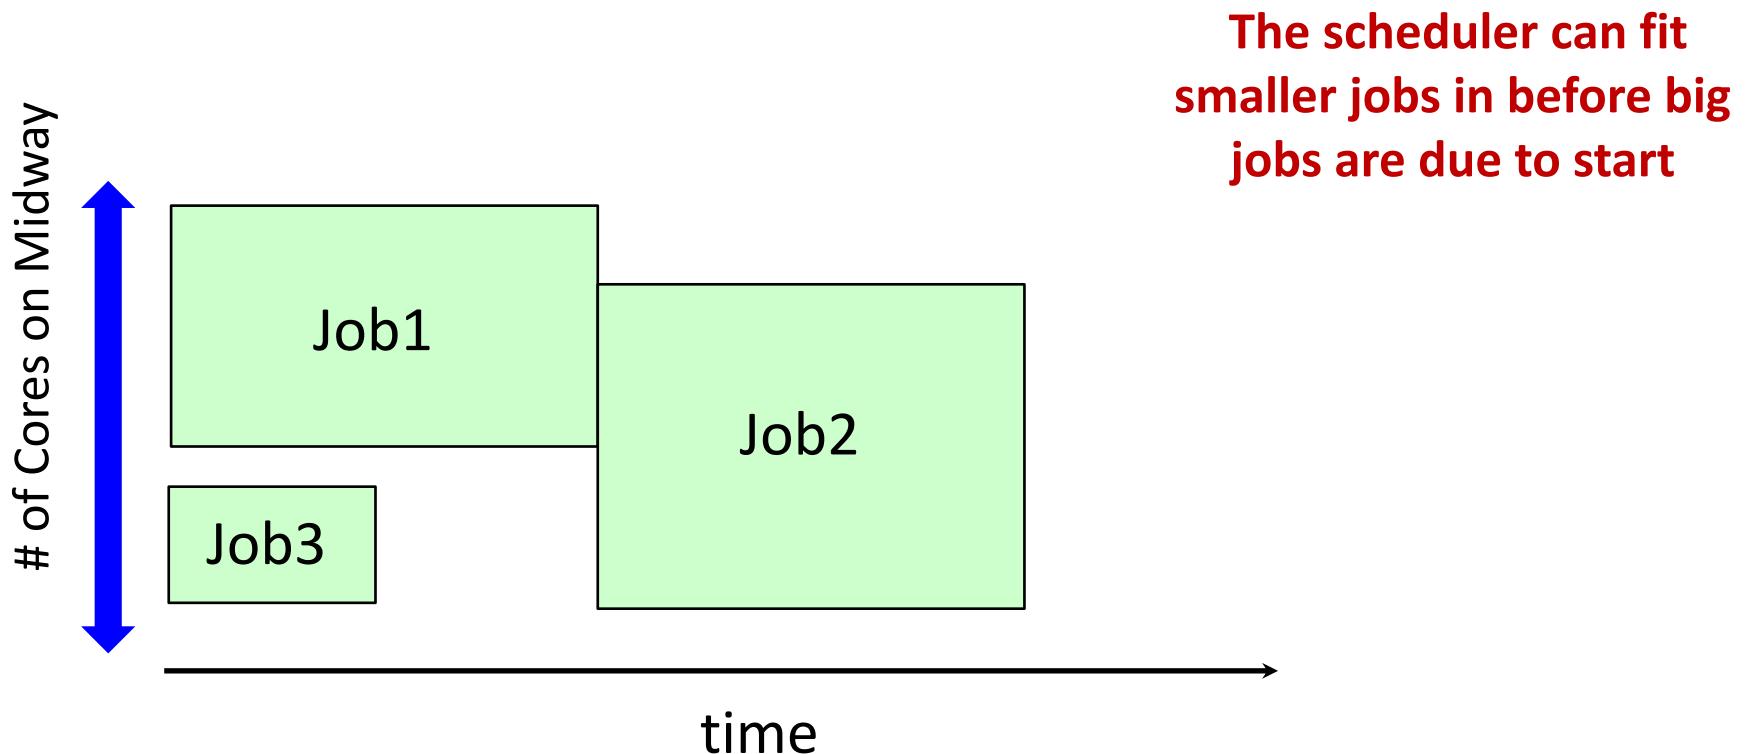
time

the x axis is time

# Job Priority

## Backfill

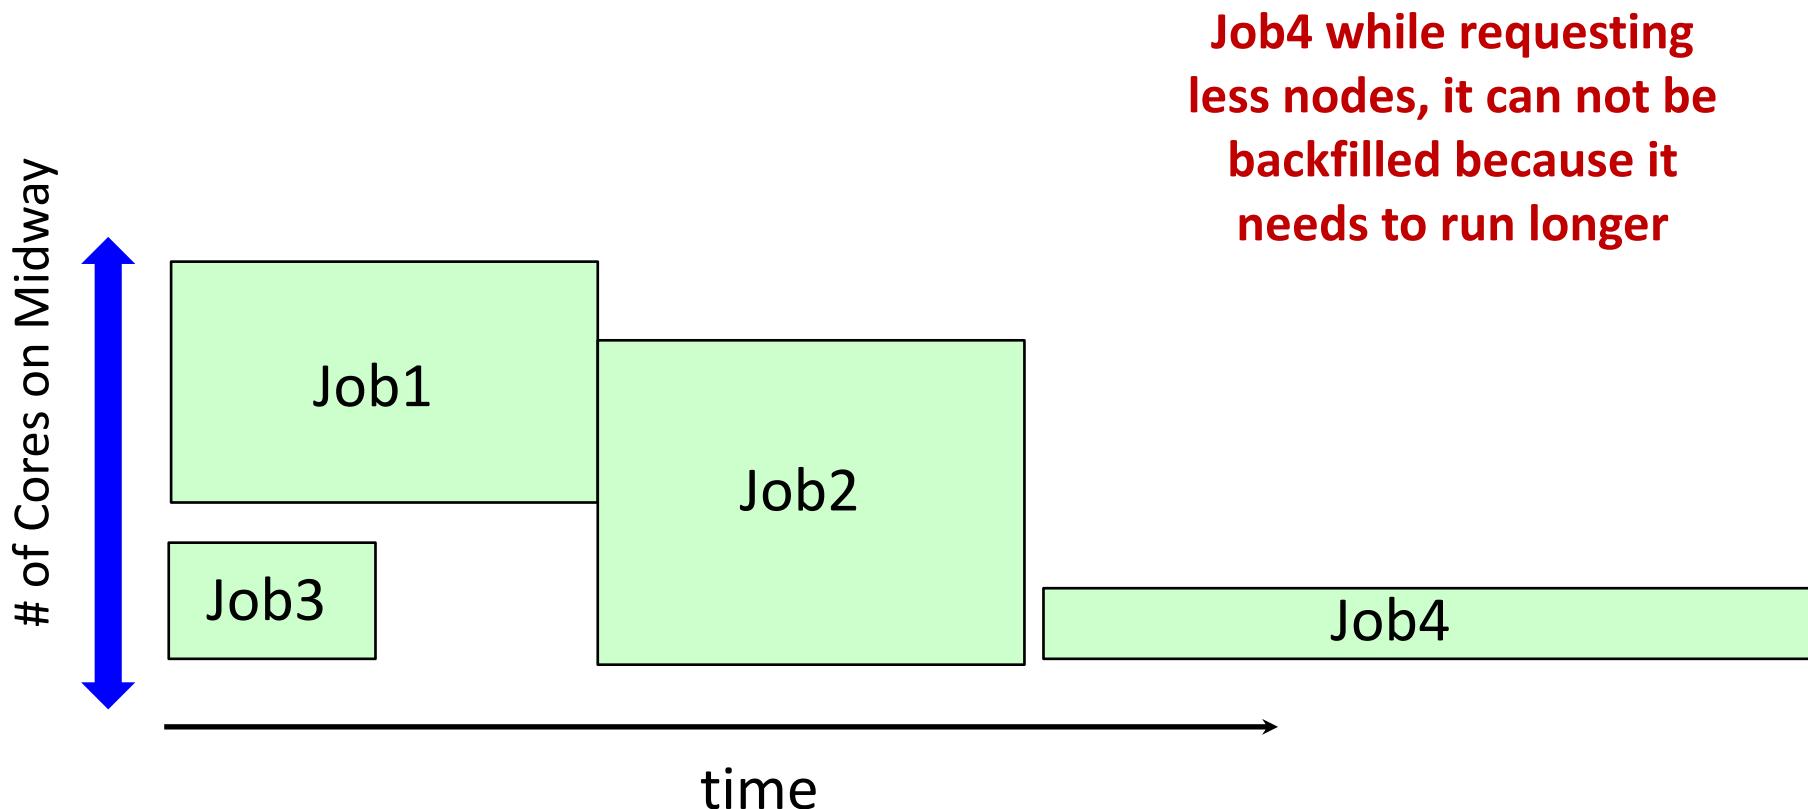
Why has someone else's job started before mine?



# Job Priority

## Backfill

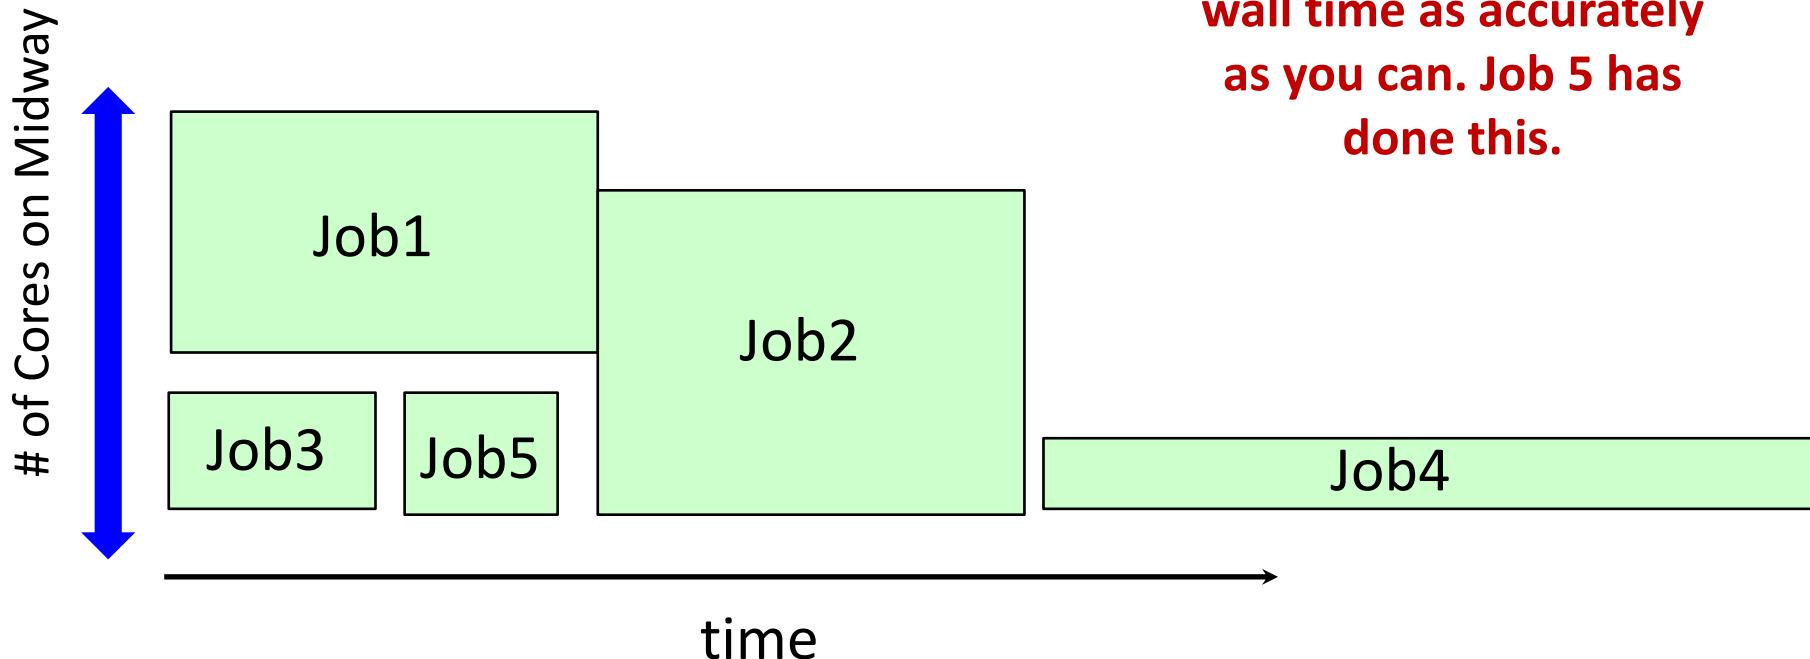
Why has someone else's job started before mine?



# Job Priority

## Backfill

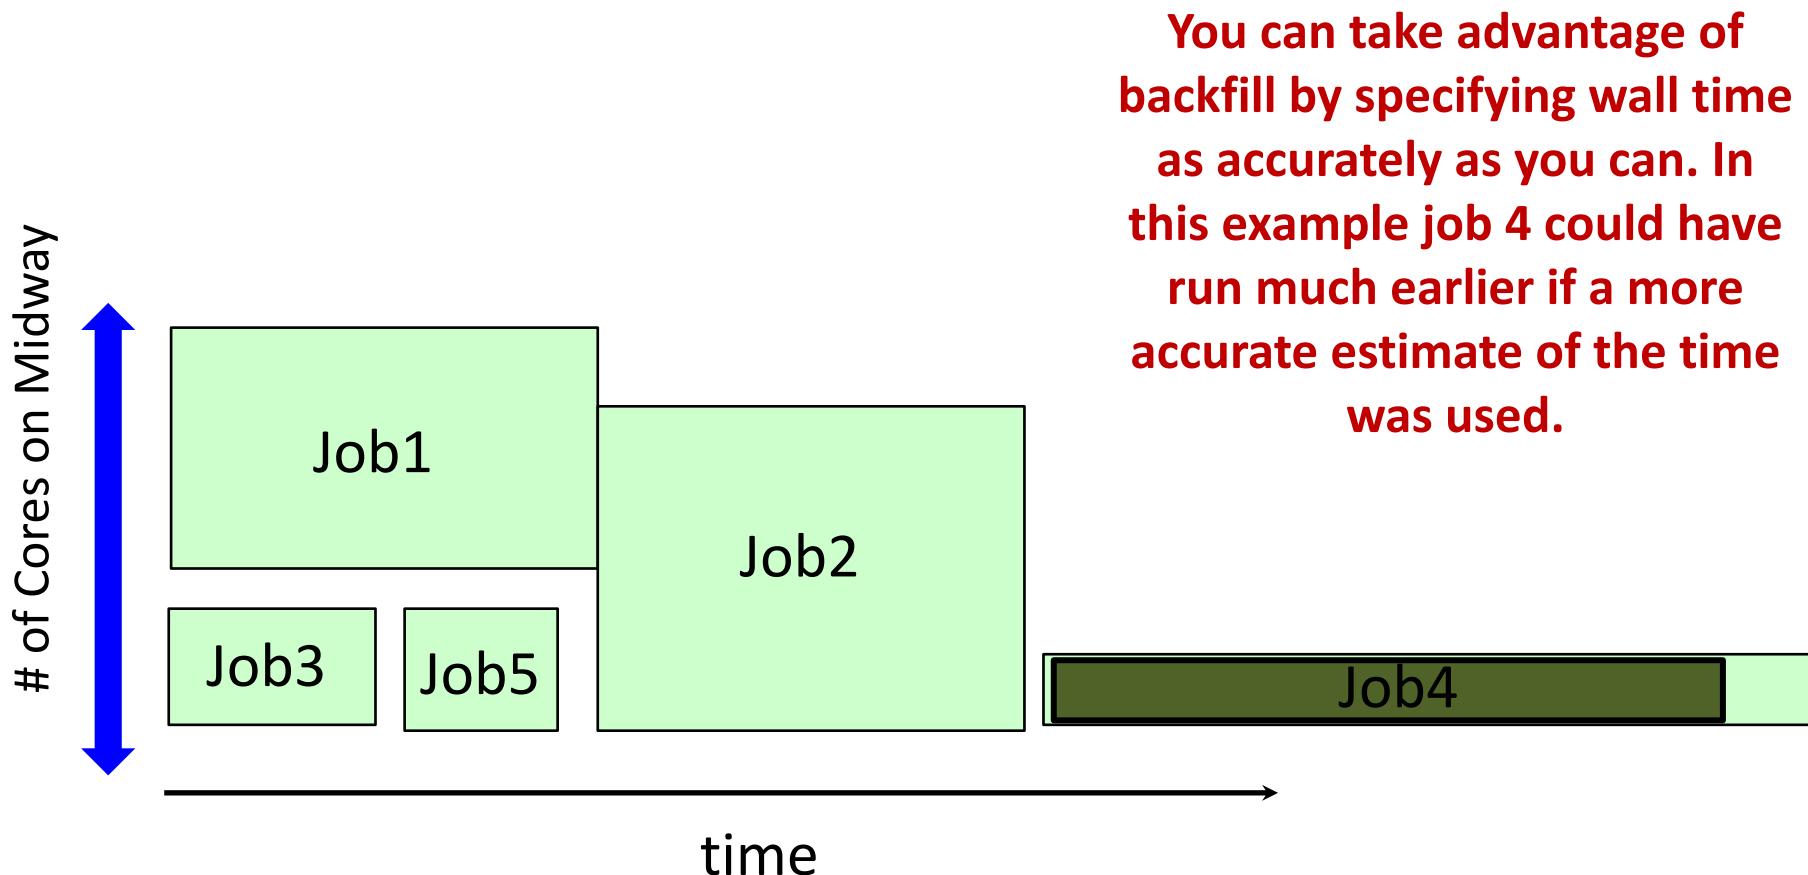
Why has someone else's job started before mine?



# Job Priority

## Backfill

Why has someone else's job started before mine?



# Job submission and monitoring

## SLURM Commands

| Command                                         | Description                                        |
|-------------------------------------------------|----------------------------------------------------|
| <code>sbatch script.sbatch</code>               | Submits <code>script.sbatch</code> job script      |
| <code>squeue -u \$USER or myq</code>            | Reports the status of your jobs                    |
| <code>sacct -u \$USER</code>                    | Displays accounting data for your job(s)           |
| <code>scancel jobid</code>                      | Cancels a running job or removes it from the queue |
| <code>scontrol show job jobid or jobinfo</code> | Displays details of a running job                  |



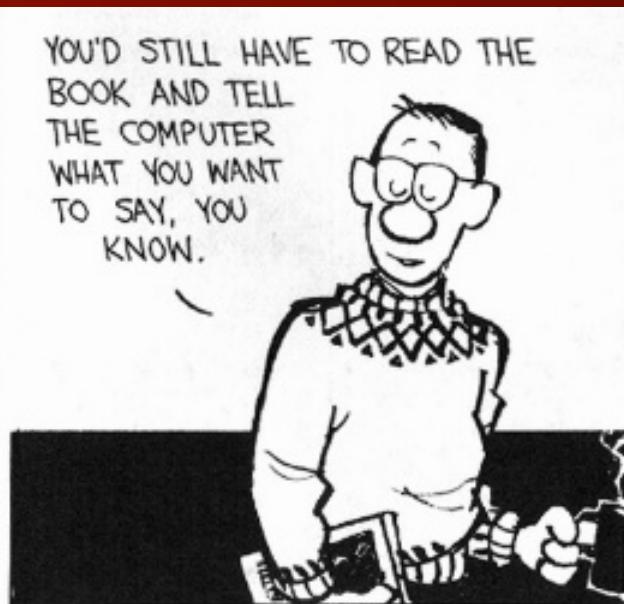
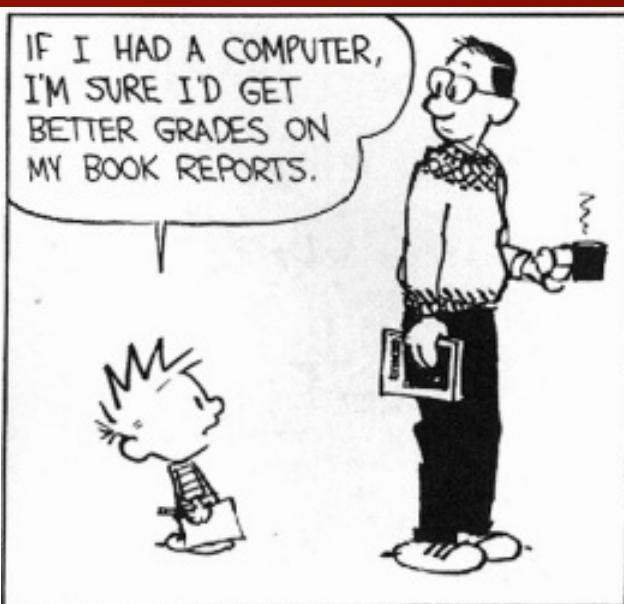
THE UNIVERSITY OF  
CHICAGO

Office of Research and  
National Laboratories  
Research Computing Center

# Recommended online resources

- User guide on running jobs on Midway
  - <https://rcc.uchicago.edu/docs/running-jobs/index.html>
- SLURM documentation
  - <https://slurm.schedmd.com/sbatch.html>
- SLURM Cheat Sheet
  - <https://slurm.schedmd.com/pdfs/summary.pdf>

# RCC Help



- Email: [help@rcc.uchicago.edu](mailto:help@rcc.uchicago.edu)
- Web: [rcc.uchicago.edu](http://rcc.uchicago.edu)
- Phone: 773-795-2667
- Walk-in: Regenstein Library, suite 216



THE UNIVERSITY OF  
CHICAGO

Office of Research and  
National Laboratories  
Research Computing Center

END OF SLIDE SHOW